

Principle of Database Systems - Project 1: Buffer Manager

18/03/2018

Southern University of Denmark

Authors: Sebastian Eklund Larsen & Kasper Kromann Nissen

Mails: selar16@student.sdu.dk, kanis16@student.sdu.dk.

Overall status

The program was implemented by implementing a buffer manager and a clock replacement policy. After the required methods were implemented, the program was tested with BMTest, and it passed successfully.

In order to implement the clock replacement policy, we created our own algorithm in the method `pickVictim`, and inserted state changes in the methods which get called by the `buffermanager`.

The other major component, namely the buffer manager itself, was implemented by the completion of the most prominent of its methods:

- `pinPage`
- `unpinPage`
- `freePage`
- `flushPage`

We implemented the methods according to their descriptions in the skeleton project. We also made use of our understanding we had gathered from reading the pages in book.

All of the methods we implemented were checked through by the both of us in a step-by-step fashion.

File descriptions

In order to implement a functioning `buffermanager`, we made changes to the following two files in the `skeleton-project` folder

Clock.java:

Implementation of a clock replacement policy.

BufMgr.java:

Implementation of the missing methods, such as `pinPage` and `flushPage`.

Division of labor

We distributed the work by having each member work on implementing a method, and once done with said method start on implementing the next unimplemented method. We encountered problems with the `pinPage` method which meant that one of us was stuck on it, and only after having both members look through the program were we able to fix it.

Sebastian Larsen:

`Clock`, `bufmgr.freePage`, `bufmgr.getNumUnpinned`, `bufmgr.flushAllPages`, `bufmgr.flushPage`.

Kasper Nissen:

bufmgr.getNumBuffers, bufmgr.pinPage, bufmgr.unpinPage, bufmgr.pinPageFound,
bufmgr.pinPageSetup.

Besides the distribution of methods by the author above, an author tag has been added to the comments of each of the implemented methods in the source code files.

Appendix

Test Output

The code was tested by running BMTest, giving the following output:

```
Running buffer manager tests...
```

```
Test 1 does a simple test of normal buffer manager operations:
```

- Allocate a bunch of new pages
- Write something on each one
- Read that something back from each one
(because we're buffering, this is where most of the writes happen)
- Free the pages again

```
Test 1 completed successfully.
```

```
Test 2 exercises some illegal buffer manager operations:
```

- Try to pin more pages than there are frames
--> Failed as expected

- Try to free a doubly-pinned page
--> Failed as expected

- Try to unpin a page not in the buffer pool
--> Failed as expected

```
Test 2 completed successfully.
```

```
Test 3 exercises some of the internals of the buffer manager
```

- Allocate and dirty some new pages, one at a time, and leave some pinned
- Read the pages

```
Test 3 completed successfully.
```

```
All buffer manager tests completed successfully!
```

Code:

BufMgr.java

```
1 package bufmgr;
2
3 import java.util.HashMap;
4
5 /**
6  * <h3>Minibase Buffer Manager</h3> The buffer manager reads disk pages into a
7  * main memory page as needed. The collection of main memory pages (called
8  * frames) used by the buffer manager for this purpose is called the buffer
9  * pool. This is just an array of Page objects. The buffer manager is used by
10 * access methods, heap files, and relational operators to read, write,
11 * allocate, and de-allocate pages.
12 */
13 @SuppressWarnings("unused")
14 public class BufMgr implements GlobalConst {
15
16     /**
17      * Actual pool of pages (can be viewed as an array of byte arrays).
18      */
19     protected Page[] bufpool;
20
21     /**
22      * Array of descriptors, each containing the pin count, dirty status, etc.
23      */
24     protected FrameDesc[] frametab;
25
26     /**
27      * Maps current page numbers to frames; used for efficient lookups.
28      */
29     protected HashMap<Integer, FrameDesc> pagemap;
30
31     /**
32      * The replacement policy to use.
33      */
34     protected Replacer replacer;
35
36     /**
37      * Constructs a buffer manager with the given settings.
38      *
39      * @param numbufs: number of pages in the buffer pool
40      */
41     public BufMgr(int numbufs) {
42         // initialize the buffer pool and frame table
43         bufpool = new Page[numbufs];
44         frametab = new FrameDesc[numbufs];
45         for (int i = 0; i < numbufs; i++) {
46             bufpool[i] = new Page();
47             frametab[i] = new FrameDesc(i);
48         }
49
50         // initialize the specialized page map and replacer
51         pagemap = new HashMap<Integer, FrameDesc>(numbufs);
52         replacer = new Clock(this);
53     }
54
55     /**
56      * Allocates a set of new pages, and pins the first one in an appropriate
57      * frame in the buffer pool.
58      */
59 }
```

BufMgr.java

```
63  *
64  * @param firstpg holds the contents of the first page
65  * @param run_size number of new pages to allocate
66  * @return page id of the first new page
67  * @throws IllegalArgumentException if PIN_MEMCPY and the page is pinned
68  * @throws IllegalStateException if all pages are pinned (i.e. pool
69  * exceeded)
70  */
71  public PageId newPage(Page firstpg, int run_size) {
72      // allocate the run
73      PageId firstid = Minibase.DiskManager.allocate_page(run_size);
74
75      // try to pin the first page
76      try {
77          pinPage(firstid, firstpg, PIN_MEMCPY);
78      } catch (RuntimeException exc) {
79          // roll back because pin failed
80          for (int i = 0; i < run_size; i++) {
81              firstid.pid += 1;
82              Minibase.DiskManager.deallocate_page(firstid);
83          }
84          // re-throw the exception
85          throw exc;
86      }
87      // notify the replacer and return the first new page id
88      replacer.newPage(pagemap.get(firstid.pid));
89      return firstid;
90  }
91
92  /**
93  *
94  * @author Sebastian Larsen
95  *
96  * Deallocates a single page from disk, freeing it from the pool if needed.
97  * Call Minibase.DiskManager.deallocate_page(pageno) to deallocate the page
98  * before return.
99  *
100  * @param pageno identifies the page to remove
101  * @throws IllegalArgumentException if the page is pinned
102  */
103  public void freePage(PageId pageno) throws IllegalArgumentException {
104      FrameDesc fdesc = pagemap.get(pageno.pid);
105      if (fdesc != null) {
106          if (fdesc.pincnt > 0) {
107              throw new IllegalArgumentException();
108          }
109          fdesc.pageno.pid = INVALID_PAGEID;
110          pagemap.remove(pageno.pid);
111          replacer.freePage(fdesc);
112      }
113      Minibase.DiskManager.deallocate_page(pageno);
114  }
115
116  /**
117  * @author Kasper Nissen
118  *
119  */
```

BufMgr.java

```
120 * Pins a disk page into the buffer pool. If the page is already pinned,
121 * this simply increments the pin count. Otherwise, this selects another
122 * page in the pool to replace, flushing the replaced page to disk if it is
123 * dirty.
124 *
125 * (If one needs to copy the page from the memory instead of reading from
126 * the disk, one should set skipRead to PIN_MEMCPY. In this case, the page
127 * shouldn't be in the buffer pool. Throw an IllegalArgumentException if so.
128 * )
129 *
130 *
131 * @param pageno identifies the page to pin
132 * @param page if skipread == PIN_MEMCPY, works as an input param,
133 * holding the contents to be read into the buffer pool if skipread ==
134 * PIN_DISKIO, works as an output param, holding the contents of the pinned
135 * page read from the disk
136 * @param skipRead PIN_MEMCPY(true) (copy the input page to the buffer
137 * pool); PIN_DISKIO(false) (read the page from disk)
138 * @throws IllegalArgumentException if PIN_MEMCPY and the page is pinned
139 * @throws IllegalStateException if all pages are pinned (i.e. pool
140 * exceeded)
141 */
142 public void pinPage(PageId pageno, Page page, boolean skipRead) {
143
144     // attempt to retrieve the FrameDesc from the pagemap
145     FrameDesc desc = pagemap.get(pageno.pid);
146
147     // if successful
148     if (desc != null) {
149         // increases the pincnt on the found page,
150         // as well as ensuring the other necessary methods are called
151         pinPageFound(page, skipRead, desc);
152         replacer.pinPage(desc);
153         return;
154     } else {
155         // find the next page to replace
156         int victim = replacer.pickVictim();
157
158         // if there is no replaceable pages
159         if (victim == -1) {
160             throw new IllegalStateException();
161         }
162
163         // replace the victim page with the new page
164         desc = frametab[victim];
165
166         // if a valid page is found at the victim index // TODO
167         if (desc.pageno.pid != INVALID_PAGEID) {
168             // remove the valid page from the pagemap
169             pagemap.remove(desc.pageno.pid);
170
171             // if dirty
172             if (desc.dirty) // write the page to disk
173             {
174                 Minibase.DiskManager.write_page(desc.pageno, bufpool[victim]);
175             }
176         }
177     }
178 }
```

BufMgr.java

```
177         pinPageSetup(pageno, page, skipRead, desc, victim);
178         replacer.pinPage(desc);
179     }
180 }
181
182 /**
183  * @author Kasper Nissen
184  *
185  * If a page already in the pagemap is to be pinned, this method will be
186  * called. It ensures that the pincnt is incremented and that the correct
187  * replacer flags are set.
188  *
189  * @param page if skipread == PIN_MEMCPY, works as as an input param,
190  * holding the contents to be read into the buffer pool if skipread ==
191  * PIN_DISKIO, works as an output param, holding the contents of the pinned
192  * page read from the disk
193  * @param skipRead PIN_MEMCPY(true) (copy the input page to the buffer
194  * pool); PIN_DISKIO(false) (read the page from disk)
195  * @param desc the FrameDesc object that holds the FrameDesc tied to the
196  * page found
197  */
198 private void pinPageFound(Page page, boolean skipRead, FrameDesc desc) {
199     if (skipRead == PIN_MEMCPY) {
200         throw new IllegalArgumentException("invalid argument, birch");
201     }
202     // pins the page
203     desc.pincnt++;
204     page.setPage(bufpool[desc.index]);
205 }
206
207 /**
208  * @author Kasper Nissen
209  *
210  * ** used in pinPage ** ensures that page is initialized with the right
211  * data, according to the value of skipRead.
212  *
213  * @param pageno identifies the page to pin
214  * @param page if skipread == PIN_MEMCPY, works as as an input param,
215  * holding the contents to be read into the buffer pool if skipread ==
216  * PIN_DISKIO, works as an output param, holding the contents of the pinned
217  * page read from the disk
218  * @param skipRead PIN_MEMCPY(true) (copy the input page to the buffer
219  * pool); PIN_DISKIO(false) (read the page from disk)
220  * @param desc the FrameDesc object that holds the FrameDesc tied to the
221  * page found
222  * @param victim the page to be replaced, determined by the replacement
223  * policy in use
224  */
225 private void pinPageSetup(PageId pageno, Page page, boolean skipRead, FrameDesc desc, int
226 victim) {
227     // if skipRead == PIN_MEMCPY
228     if (skipRead) {
229         // copy from memory
230         bufpool[victim].copyPage(page);
231     } // if skipRead == PIN_DISKIO
232 }
```


BufMgr.java

```
233     else {
234         // read from disk
235         Minibase.DiskManager.read_page(pageno, bufpool[victim]);
236     }
237
238     desc.pincnt = 1;
239     page.setPage(bufpool[victim]);
240     pagemap.put(pageno.pid, desc);
241     desc.pageno.pid = pageno.pid;
242 }
243
244 /**
245  * @author Kasper Nissen
246  *
247  * Unpins a disk page from the buffer pool, decreasing its pin count.
248  *
249  * @param pageno identifies the page to unpin
250  * @param dirty UNPIN_DIRTY if the page was modified, UNPIN_CLEAN otherwise
251  * @throws IllegalArgumentException if the page is not present or not pinned
252  */
253 public void unpinPage(PageId pageno, boolean dirty) throws IllegalArgumentException {
254
255     if (!pagemap.containsKey(pageno.getPID())) {
256         throw new IllegalArgumentException();
257     }
258
259     FrameDesc desc = pagemap.get(pageno.getPID());
260
261     if (desc.pincnt == 0) {
262         throw new IllegalArgumentException();
263     }
264
265     desc.pincnt--;
266
267     if (dirty == UNPIN_DIRTY) {
268         desc.dirty = true;
269     }
270     replacer.unpinPage(desc);
271 }
272
273 /**
274  *
275  * @author Sebastian Larsen
276  *
277  * Immediately writes a page in the buffer pool to disk, if dirty.
278  *
279  */
280 public void flushPage(PageId pageno) {
281
282     FrameDesc fdesc = pagemap.get(pageno.pid);
283
284     if (fdesc.dirty) {
285         Minibase.DiskManager.write_page(pageno, bufpool[fdesc.index]);
286         fdesc.dirty = false;
287     }
288 }
289
```

BufMgr.java

```
290  /**
291  *
292  * @author Sebastian Larsen
293  *
294  * Immediately writes all dirty pages in the buffer pool to disk, skipping
295  * the pages which have INVALID_PAGEID
296  */
297  public void flushAllPages() {
298      for (int i = 0; i < frametab.length; i++) {
299          if (frametab[i].pageno.pid != INVALID_PAGEID) {
300              flushPage(frametab[i].pageno);
301          }
302      }
303  }
304
305  /**
306  * @author Kasper Nissen
307  *
308  * Gets the total number of buffer frames.
309  */
310  public int getNumBuffers() {
311      return frametab.length;
312  }
313
314  /**
315  *
316  * @author Sebastian Larsen
317  *
318  * Gets the total number of unpinned buffer frames.
319  */
320  public int getNumUnpinned() {
321      int unpin_count = 0;
322      for (int i = 0; i < frametab.length; i++) {
323          if (frametab[i].pincnt == 0) {
324              unpin_count++;
325          }
326      }
327      return unpin_count;
328  }
329
330 } // public class BufMgr implements GlobalConst
331
```

Clock.java

```
1 package bufmgr;
2 /**
3  *
4  * @author Sebastian Larsen
5  *
6  * The clock replacement policy looks at the frametable as a clock
7  * and loops through the indices evaluating them based on their state.
8  *
9  * The state of an index can be one of the following 3
10 * 0 = available : the index can be replaced
11 * 1 = pinned : the index cannot be replaced
12 * 2 = prevpinned : should be set to available instead of being replaced when picked.
13 */
14 public class Clock extends Replacer{
15
16     private int current;
17
18     protected Clock(BufMgr bufmgr) {
19         super(bufmgr);
20         current = -1;
21     }
22
23     @Override
24     public void newPage(FrameDesc fdesc) {
25         // TODO Auto-generated method stub
26     }
27
28
29     @Override
30     public void freePage(FrameDesc fdesc) {
31
32         fdesc.state=0;
33     }
34
35
36     @Override
37     public void pinPage(FrameDesc fdesc) {
38         fdesc.state=1;
39     }
40
41
42     @Override
43     public void unpinPage(FrameDesc fdesc) {
44         if(fdesc.pincnt == 0){
45             fdesc.state=2;
46         }
47     }
48
49
50     @Override
51     /**
52     *
53     * pickVictim tries to pick a frame which can be replaced.
54     * Loops through the frametable upto 2 times looking for a frame to be replaced.
55     * For each frame it reaches it checks if its state is available(0), in which the index
56     is returned,
57     * or if the state is prevpinned(2), in which case it is set to available.
```

Clock.java

```
57      * If no index is found in two cycles through the frametable, it means
58      * that every frame is unavailable, and -1 is returned
59      *
60      *
61      */
62      public int pickVictim() {
63          int count = 0;
64          while (count < frametab.length * 2) {
65              current = (current + 1) % frametab.length;
66
67              if (frametab[current].state == 0) {
68                  return current;
69              } else if (frametab[current].state == 2) {
70                  frametab[current].state = 0;
71              }
72              count++;
73          }
74          return -1;
75      }
76 }
77
```